

Efficiency through Uncertainty: Scalable Formal Synthesis for Stochastic Hybrid Systems

Nathalie Cauchi¹ Luca Laurenti¹ Morteza Lahijanian²,
Alessandro Abate¹ Marta Kwiatkowska¹ Luca Cardelli^{1,3}


HSCC 2019

¹University of Oxford

²University of Colorado Boulder

³Microsoft Research at Cambridge

Introduction



Safety critical systems with a multitude of requirements

System dynamics:

- complex continuous dynamics
- discrete modes
- uncertainty


Specifications:

- can be formally defined

stochastic hybrid systems + formal methods

Classical approach



- **abstract** continuous variables to **finite** Markov Processes
- errors between original and abstract model (treated as separate parameter)



Introduction

Classical approach

- **abstract** continuous variables to **finite** Markov Processes
- errors between original and abstract model (treated as separate parameter)
- generate **conservative** error bounds
 - error increases linearly with time
 - state-space explosion due to error explosion
- limited to finite time properties
- synthesis over abstract model (classical MDP synthesis)



Problem formulation: stochastic hybrid systems

A (discrete-time) linear stochastic hybrid system (SHS) is a tuple

$$\mathcal{H} = (\mathcal{A}, F, G, \Upsilon, L)$$




discrete modes
 $\mathcal{A} = \{a_0, \dots, a_{|\mathcal{A}|}\}$



Problem formulation: stochastic hybrid systems


A (discrete-time) linear stochastic hybrid system (SHS) is a tuple

$$\mathcal{H} = (\mathcal{A}, F, G, \Upsilon, L)$$



drift
 $F = \{F(a) \in \mathbb{R}^{m \times m} \mid a \in \mathcal{A}\}$


diffusion
 $G = \{G(a) \in \mathbb{R}^{m \times r} \mid a \in \mathcal{A}\}$




Problem formulation: stochastic hybrid systems

A (discrete-time) linear stochastic hybrid system (SHS) is a tuple

$$\mathcal{H} = (\mathcal{A}, F, G, \Upsilon, L)$$




$\Upsilon = \{p_1, \dots, p_n\}$
atomic propositions
 $L : S \rightarrow 2^\Upsilon$
labeling function



Problem formulation: stochastic hybrid systems

A (discrete-time) linear stochastic hybrid system (SHS) is a tuple


$$\mathcal{H} = (\mathcal{A}, F, G, \Upsilon, L)$$



Problem formulation: stochastic hybrid systems

A (discrete-time) linear stochastic hybrid system (SHS) is a tuple

$$\mathcal{H} = (\mathcal{A}, F, G, \Upsilon, L)$$



stochastic process

- a hybrid state of \mathcal{H} is a pair $s = (a, x) \in S$
- evolution of \mathcal{H} for $k \in \mathbb{Z}_{\geq 0}$ is a stochastic process $s(k) = (a(k), x(k)) \in S$

Problem formulation: stochastic hybrid systems

switching strategy

a function that assigns a discrete mode $a \in \mathcal{A}$ to a finite path $\omega_{\mathcal{H}}$ of the process s

Problem formulation: stochastic hybrid systems

switching strategy

a function that assigns a discrete mode $a \in \mathcal{A}$ to a finite path $\omega_{\mathcal{H}}$ of the process s

transition kernel

for any measurable set $B \subseteq \mathbb{R}^m$, $x \in \mathbb{R}^m$, and $a \in \mathcal{A}$

$$T(B \mid x, a) = \int_B \mathcal{N}(t \mid F(a)x, G(a)^T \text{Cov}_w G(a)) dt$$

Problem formulation: stochastic hybrid systems

switching strategy

a function that assigns a discrete mode $a \in \mathcal{A}$ to a finite path $\omega_{\mathcal{H}}$ of the process s

transition kernel

for any measurable set $B \subseteq \mathbb{R}^m$, $x \in \mathbb{R}^m$, and $a \in \mathcal{A}$

$$T(B \mid x, a) = \int_B \mathcal{N}(t \mid F(a)x, G(a)^T \text{Cov}_w G(a)) dt$$

temporal logic specifications

- deals with complex formal properties with boolean and temporal constraints (bounded and unbounded)

Problem statement

Given:

1. SHS \mathcal{H}
2. a compact set X
3. property expressed as a formula φ defined over regions of X
4. φ also requires system not to leave X

Problem statement

Given:

1. SHS \mathcal{H}
2. a compact set X
3. property expressed as a formula φ defined over regions of X
4. φ also requires system not to leave X

Find:

1. a **switching strategy** that **maximizes** the probability of satisfying φ for all initial states $s_0 \in \mathcal{A} \times X$


Proposed approach

We construct abstract model that captures all behaviours of SHS with to X and regions of interest

Proposed approach

We construct abstract model that captures all behaviours of SHS with to X and regions of interest


1. **discretize** set X according to dynamics of each mode



Proposed approach

We construct abstract model that captures all behaviours of SHS with to X and regions of interest


1. **discretize** set X according to dynamics of each mode
2. **quantify** error of abstraction and represent it as uncertainty in the abstraction



Proposed approach

We construct abstract model that captures all behaviours of SHS with to X and regions of interest


1. discretize set X
2. quantify error of abstraction and represent it as uncertainty
3. synthesise optimal strategy via stochastic games
4. strategy can be mapped back to SHS



Preliminaries: models

An IMDP is a tuple $\mathcal{I} = (Q, A, \check{P}, \hat{P})$

- Q is a finite set of states
- A is a finite set of actions
- $\check{P}(q, a, q')$ defines the lower bound of the transition probability
- $\hat{P}(q, a, q')$ defines the upper bound of the transition probability
- Υ is a finite set of atomic propositions
- $L : Q \rightarrow 2^\Upsilon$ is a labeling function




An IMDP is a tuple $\mathcal{I} = (Q, A, \check{P}, \hat{P})$

- a **feasible distribution** reachable from q by a if

$$\check{P}(q, a, q') \leq \gamma_q^a(q') \leq \hat{P}(q, a, q')$$


for each state $q' \in Q$



SHS abstractions

We abstract the SHS to an IMDP



$$\mathcal{I} = (Q, A, \check{P}, \hat{P}, \Upsilon, L)$$



SHS abstractions

We abstract the SHS to an IMDP

$$\mathcal{I} = (Q, A, \check{P}, \hat{P}, \Upsilon, L)$$



SHS abstractions

We abstract the SHS to an IMDP

$$\mathcal{I} = (\mathcal{Q}, A, \check{P}, \hat{P}, \Upsilon, L)$$

$q_1^{a_1}$	$q_2^{a_1}$	$q_3^{a_1}$	$q_4^{a_1}$
$q_5^{a_1}$	$q_6^{a_1}$	$q_7^{a_1}$	$q_8^{a_1}$
$q_9^{a_1}$	$q_{10}^{a_1}$	$q_{11}^{a_1}$	$q_{12}^{a_1}$

$$\mathcal{Q}^{a_1} = \{q_1^{a_1}, \dots, q_{12}^{a_1}\}$$

$q_1^{a_2}$	$q_2^{a_2}$	$q_3^{a_2}$	$q_4^{a_2}$
$q_5^{a_2}$	$q_6^{a_2}$	$q_7^{a_2}$	$q_8^{a_2}$
$q_9^{a_2}$	$q_{10}^{a_2}$	$q_{11}^{a_2}$	$q_{12}^{a_2}$

$$\mathcal{Q}^{a_2} = \{q_1^{a_2}, \dots, q_{12}^{a_2}\}$$

$$\mathcal{Q} = \{\mathcal{Q}^{a_1} \cup \mathcal{Q}^{a_2} \cup \{q_u\}\}$$

We abstract the SHS to an IMDP

$$\mathcal{I} = (Q, A, \check{P}, \hat{P}, \Upsilon, L)$$

- A are the set of **modes** of \mathcal{H} : $A(q) = A \ \forall q \in Q$

We abstract the SHS to an IMDP

$$\mathcal{I} = (Q, A, \check{P}, \hat{P}, \Upsilon, L)$$

- A are the set of **modes** of \mathcal{H} : $A(q) = A \forall q \in Q$
- one-step transition probability : $T(q|x, a)$
- **but** $q \in Q$ correspond to **regions** in \mathcal{H}
 - **range** of feasible transition probabilities to region q
 - **bound** feasible transitions to get \check{P}, \hat{P}

We abstract the SHS to an IMDP

$$\mathcal{I} = (Q, A, \check{P}, \hat{P}, \Upsilon, L)$$

- A are the set of **modes** of \mathcal{H} : $A(q) = A \forall q \in Q$
- one-step transition probability : $T(q|x, a)$

$$\gamma_{q_i}^a(q_j) \leq \min_{x \in q_i} T(q_j|x, a)$$

$$\gamma_{q_i}^a(q_j) \geq \max_{x \in q_i} T(q_j|x, a)$$

Computation of imdp


How do we efficiently compute

$$\min_{x \in q_i} T(q_j \mid x, a), \max_{x \in q_i} T(q_j \mid x, a)?$$

Computation of imdp

How do we efficiently compute

$$\min_{x \in q_i} T(q_j \mid x, a), \quad \max_{x \in q_i} T(q_j \mid x, a)?$$





For process x in mode $a \in \mathcal{A}$

Transformation

$$\mathcal{T}_a = \Lambda_a^{-\frac{1}{2}} V_a^T$$

$\Lambda_a = V_a^T \text{Cov}_x(a) V_a$






Hyper rectangles

$$[v_l^{(1)}, v_u^{(1)}] \times \cdots \times [v_l^{(m)}, v_u^{(m)}]$$

Computation of imdp

How do we efficiently compute

$$\min_{x \in q_i} T(q_j | x, a), \quad \max_{x \in q_i} T(q_j | x, a)?$$




For process x in mode $a \in \mathcal{A}$

Transformation

$$\mathcal{T}_a = \Lambda_a^{-\frac{1}{2}} V_a^T$$

$$\Lambda_a = V_a^T \text{Cov}_x(a) V_a$$



Analytical solution


$$[v_l^{(1)}, v_u^{(1)}] \times \dots \times [v_l^{(m)}, v_u^{(m)}]$$

$$T(q_j | x, a) = \frac{1}{2^m} \prod_{n=1}^m \left(\text{erf}\left(\frac{y^{(n)} - v_l^{(n)}}{\sqrt{2}}\right) - \text{erf}\left(\frac{y^{(n)} - v_u^{(n)}}{\sqrt{2}}\right) \right)$$

Computation of imdp

How do we efficiently compute

$$\min_{x \in q_i} T(q_j \mid x, a), \quad \max_{x \in q_i} T(q_j \mid x, a)?$$



$$\min_{x \in q_i} T(q_j \mid x, a) = \min_{y \in \text{Post}(q'_j, T_a)} f(y),$$

$$\max_{x \in q_i} T(q_j \mid x, a) = \max_{y \in \text{Post}(q'_j, T_a)} f(y).$$

$$f(y) = \frac{1}{2^m} \prod_{n=1}^m \left(\text{erf}\left(\frac{y^{(n)} - v_l^{(n)}}{\sqrt{2}}\right) - \text{erf}\left(\frac{y^{(n)} - v_u^{(n)}}{\sqrt{2}}\right) \right)$$

Optimisation via KKT or Gradient descent


Optimisation via KKT conditions

- solving systems of non-linear equations
- efficient and exact for low-dimensional system
- number of vertices to check grows exponentially with dimensions

Computation of imdp

Optimisation via KKT conditions


- solving systems of non-linear equations
- efficient and exact for low-dimensional system
- number of vertices to check grows exponentially with dimensions



Computation of imdp

Optimisation via KKT conditions

- solving systems of non-linear equations
- efficient and exact for low-dimensional system
- number of vertices to check grows exponentially with dimensions



Optimisation via **gradient descent** method

$$f(y) = \frac{1}{2^m} \left[\prod_{i=1}^m \left(\operatorname{erf}\left(\frac{y^{(i)} - v_l^{(i)}}{\sqrt{2}}\right) - \operatorname{erf}\left(\frac{y^{(i)} - v_u^{(i)}}{\sqrt{2}}\right) \right) \right]$$

Computation of imdp

Optimisation via **gradient descent** method

$$f(y) = -\frac{1}{2^m} \log \left[\prod_{i=1}^m \left(\operatorname{erf}\left(\frac{y^{(i)} - v_l^{(i)}}{\sqrt{2}}\right) - \operatorname{erf}\left(\frac{y^{(i)} - v_u^{(i)}}{\sqrt{2}}\right) \right) \right]$$

- $f(y)$ has the property of being **log-concave**
- can use standard **convex** optimisation techniques
- allows for **scaling** to **high** dimensions

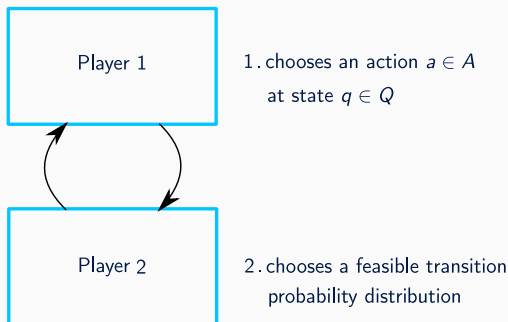
We abstract the SHS to an IMDP

$$\mathcal{I} = (Q, A, \check{P}, \hat{P}, \Upsilon, L)$$

- associate labels with corresponding region R in X
- when discretisation does not respect R
 - add extra labels (conservatively)
 - converting φ into negation normal form (NNF)
 - associate labels with negation of propositions
 - under approximate this region

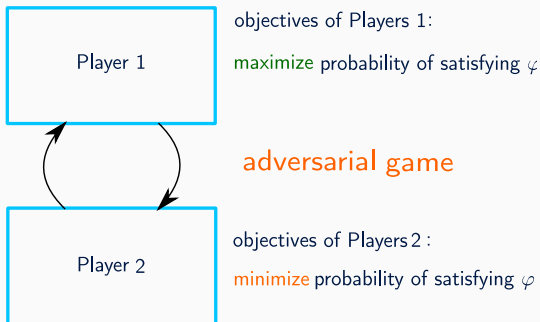
Strategy synthesis

- **uncertainties** in \mathcal{I} : **nondeterministic choice** of transition probability from one IMDP state to another under a given action
- synthesis task interpreted as a **2 $\frac{1}{2}$ player stochastic game**
- the set of actions of player 2 is **continuous**



Strategy synthesis

- **uncertainties** in \mathcal{I} : **nondeterministic choice** of transition probability from one IMDP state to another under a given action
- synthesis task interpreted as a **2 $\frac{1}{2}$ player stochastic game**
- the set of actions of player 2 is **continuous**



- implemented the **abstraction** and **synthesis** algorithms
- **test performance** using three case studies
- analysis on the **abstraction error** generated

$$\varepsilon_{max} = \max_{q \in Q} (\hat{p}(q) - \check{p}(q))$$

$\min \check{p}(q)$ and $\max \hat{p}(q)$ probabilities of satisfaction for each state q

- case studies are run on an Intel Core i7-8550U CPU at 1.80GHz \times 8 machine with 8 GB of RAM

Case study 1: formal verification

model

1 discrete mode: $A = \{a_1\}$

$$x(k+1) = \begin{pmatrix} 0.85 & 0 \\ 0 & 0.90 \end{pmatrix} x(k) + \begin{pmatrix} 0.15 & 0 \\ 0 & 0.05 \end{pmatrix} w(k)$$

Case study 1: formal verification

model

1 discrete mode: $A = \{a_1\}$

$$x(k+1) = \begin{pmatrix} 0.85 & 0 \\ 0 & 0.90 \end{pmatrix} x(k) + \begin{pmatrix} 0.15 & 0 \\ 0 & 0.05 \end{pmatrix} w(k)$$

verification of a **safety** property

$$\varphi_1 = \mathcal{G}^{\leq K} \{[-1, 1] \times [-1, 1]\}$$

comparison against state of the art tool FAUST²

Case study 1: formal verification

results:

$$\varphi_1 = \mathcal{G}^{\leq K=2} \{[-1, 1] \times [-1, 1]\}$$

Tool Method	Impl. Platform	$ \bar{Q} $ (states)	Time taken (secs)	Error ε_{\max}
IMDP (KKT)	MATLAB	361	19.789	0.211
FAUST ²	MATLAB	361	108.265	1.000
IMDP (KKT)	MATLAB	625	145.563	0.163
FAUST ²	MATLAB	625	285.795	1.000
IMDP (KKT)	MATLAB	2601	28127.256	0.082
FAUST ²	MATLAB	2601	5274.578	0.995
IMDP (KKT)	MATLAB	3721	Time out ^a	-
FAUST ²	MATLAB	3721	11285.313	0.832

^a9 hours+ and no solution

Case study 1: formal verification

results:

$$\varphi_1 = \mathcal{G}^{\leq K=2} \{[-1, 1] \times [-1, 1]\}$$


Tool Method	Impl. Platform	$ \bar{Q} $ (states)	Time taken (secs)	Error ε_{\max}
IMDP (KKT)	MATLAB	361	19.789	0.211
FAUST ²	MATLAB	361	108.265	1.000
IMDP (KKT)	MATLAB	625	145.563	0.163
FAUST ²	MATLAB	625	285.795	1.000
IMDP (KKT)	MATLAB	2601	28127.256	0.082
FAUST ²	MATLAB	2601	5274.578	0.995
IMDP (KKT)	MATLAB	3721	Time out ^a	-
FAUST ²	MATLAB	3721	11285.313	0.832

Tool Method	Impl. Platform	$ \bar{Q} $ (states)	Time taken (secs)	Error ε_{\max}
IMDP (GD)	C++	361	29.003	0.211
FAUST ²	C++	361	136.71	1.000
IMDP (GD)	C++	625	117.741	0.163
FAUST ²	C++	625	302.900	1.000
IMDP (GD)	C++	2601	2939.050	0.082
FAUST ²	C++	2601	3305.490	0.995
IMDP (GD)	C++	3721	3973.28	0.068
FAUST ²	C++	3721	7537.750	0.832


^a9 hours+ and no solution

Case study 1: formal verification

results:



FAUST²





IMDP

Case study 1: formal verification

results:



$$\varphi_1 = \mathcal{G}^{\leq k}\{[-1, 1] \times [-1, 1]\}$$



Case study 1: formal verification

results:

$$\varphi_1 = \mathcal{G}^{\leq k}\{[-1, 1] \times [-1, 1]\}$$



- error **embedded** within abstraction
- performs computations according to **feasible** transition probabilities
- abstraction errors **does not** explode with time

Case study 2: strategy synthesis

model

2 discrete modes: $A = \{a_1, a_2\}$

$$x(k+1)_{a_1} = \begin{pmatrix} 0.1 & 0.9 \\ 0.8 & 0.2 \end{pmatrix} x(k)_{a_1} + \begin{pmatrix} 0.3 & 0.1 \\ 0.1 & 0.2 \end{pmatrix} w_{a_1}(k)$$

$$x(k+1)_{a_2} = \begin{pmatrix} 0.8 & 0.2 \\ 0.1 & 0.9 \end{pmatrix} x(k)_{a_2} + \begin{pmatrix} 0.2 & 0 \\ 0 & 0.1 \end{pmatrix} w_{a_2}(k)$$

Case study 2: strategy synthesis

model

2 discrete modes: $A = \{a_1, a_2\}$


$$x(k+1)_{a_1} = \begin{pmatrix} 0.1 & 0.9 \\ 0.8 & 0.2 \end{pmatrix} x(k)_{a_1} + \begin{pmatrix} 0.3 & 0.1 \\ 0.1 & 0.2 \end{pmatrix} w_{a_1}(k)$$

$$x(k+1)_{a_2} = \begin{pmatrix} 0.8 & 0.2 \\ 0.1 & 0.9 \end{pmatrix} x(k)_{a_2} + \begin{pmatrix} 0.2 & 0 \\ 0 & 0.1 \end{pmatrix} w_{a_2}(k)$$

synthesis of a maximising **switching strategy**

$$\varphi_2 = \neg \text{red} \mathcal{U} \text{green}$$


within the set $X = [-2, 2] \times [-2, 2]$




Case study 2: strategy synthesis

results:

- total number of strates: 3612
- computational time: 5434 [s]




lower bound probability for a_1




lower bound probability for a_2

Case study 2: strategy synthesis

results:



original set X with simulated trajectories



maximum error incurred in satisfying φ_2

Case study 3: scaling in continuous dimension of model

model with d continuous variables

1 discrete mode: $A = \{a_1\}$

$$x(k+1) = -0.95\mathbf{1}_d x(k) + 0.11\mathbf{1}_d w(k)$$

Case study 3: scaling in continuous dimension of model

model with d continuous variables

1 discrete mode: $A = \{a_1\}$

$$x(k+1) = -0.951_d x(k) + 0.11_d w(k)$$



verification of a **safety** property

$$\varphi_1 = \mathcal{G}^{\leq k=50}[-1, 1]^d$$

Case study 3: scaling in continuous dimension of model

results:



- current state of the art



Case study 3: scaling in continuous dimension of model

results:



- number of dimensions that can now be analysed
- adaptive refinement resulting in max error of 0.03



Case study 3: scaling in continuous dimension of model

results:

- number of dimensions that can now be analysed
- adaptive refinement resulting in max error of 0.03
- manageable state spaces
- remarkable improvement



contributions

- theoretical & computational **technique** for **analysis** of SHS
- **precise** and **compact** abstractions
- algorithms embedded within new tool **StocHy** presented at TACAS'19

Conclusion

contributions

- theoretical & computational **technique** for **analysis** of SHS
- **precise** and **compact** abstractions
- algorithms embedded within new tool **StocHy** presented at TACAS'19

future work

- synthesis for more complex and even multi-objective properties
- analysis of continuous-time SHS (**current work in progress**)

contributions

- theoretical & computational **technique** for **analysis** of SHS
- **precise** and **compact** abstractions
- algorithms embedded within new tool **StocHy** presented at TACAS'19


Thank you!

abstraction into MDP

- **define** desired abstraction error
- grid state-space (**uniform** or **adaptive**)
- compute transition probabilities via marginalisation
- hinges on computation of **Lipschitz** constants (h_s)
- N step error

$$\varepsilon = h_s \delta \mathcal{L}(A) N$$

δ max diameter of partition, $\mathcal{L}(A)$
volume of set





abstraction into MDP

- define desired abstraction error
- grid state-space (uniform or adaptive)
- compute transition probabilities via marginalisation
- hinges on computation of Lipschitz constants (h_s)
- N step error

$$\varepsilon = h_s \delta \mathcal{L}(A) N$$

δ max diameter of partition, $\mathcal{L}(A)$
volume of set





abstraction into MDP

- define desired abstraction error
- grid state-space (uniform or adaptive)
- compute transition probabilities via marginalisation
- hinges on computation of Lipschitz constants (h_s)
- N step error

$$\varepsilon = h_s \delta \mathcal{L}(A) N$$

δ max diameter of partition, $\mathcal{L}(A)$
volume of set






abstraction into MDP

- define desired abstraction error
- grid state-space (uniform or adaptive)
- compute transition probabilities via marginalisation
- hinges on computation of Lipschitz constants (h_s)
- N step error

$$\varepsilon = h_s \delta \mathcal{L}(A) N$$

δ max diameter of partition, $\mathcal{L}(A)$
volume of set



- conservatively overapproximate discretizations of $\mathcal{A} \times X$ that donot respect regions in $R = \{r_1, \dots, r_n\}$
- represent each region by its complement relative to X
- Let $r_{n+i} = X \setminus r_i$ be the complement region of r_i with respect to X . We associate to each r_{n+i} a new atomic proposition p_{n+i} for $1 \leq i \leq n$.

$$\bar{\Upsilon} = \Upsilon \cup \{p_{n+1}, \dots, p_{2n}\}$$

Then, we design $L : Q \rightarrow 2^{\bar{\Upsilon}}$ of \mathcal{I} such that

$$p_i \in L(q) \quad \Leftrightarrow \quad q \subseteq r_i$$

for all $q \in \bar{Q}$ and $0 \leq i \leq 2n$, and $L(q_u) = \emptyset$